

import the libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import sklearn
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.metrics import accuracy_score
from sklearn import svm
```

load the data

```
In [2]: data = pd.read_csv("/home/tamanna/Downloads/creditcard.csv")
data
```

```
Out[2]:
```

	Time	V1	V2	V3	V4	V5	V6	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.23
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.07
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.79
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.23
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.59
...
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.91
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.02
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.29
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.68
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.57

284807 rows × 31 columns

data preprocessing

```
In [3]: data.head()
```

Out[3]:

	Time	V1	V2	V3	V4	V5	V6	V7	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270

5 rows × 31 columns

In [4]: data.tail()

Out[4]:

	Time	V1	V2	V3	V4	V5	V6	V7	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.290	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.680	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577	

5 rows × 31 columns

In [5]: # generating descriptive statistics
data.describe()

Out[5]:

	Time	V1	V2	V3	V4	V5	V6	V7
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	1.168375e-15	3.416908e-16	-1.379537e-15	2.074095e-15	9.604095e-16	-1.379537e-15	2.074095e-15
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.380145e+00	1.415869e+00	1.380145e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137145e+01	-5.683171e+00	-1.137145e+01
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.915499e-01	-8.486401e-01	-6.915499e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.433080e-02	-1.984653e-02	-5.433080e-02
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.115499e-01	7.433413e-01	6.115499e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.480145e+00	1.687534e+01	3.480145e+00

8 rows × 31 columns

In [6]: data.shape

Out[6]: (284807, 31)

In [7]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Time        284807 non-null  float64
1   V1          284807 non-null  float64
2   V2          284807 non-null  float64
3   V3          284807 non-null  float64
4   V4          284807 non-null  float64
5   V5          284807 non-null  float64
6   V6          284807 non-null  float64
7   V7          284807 non-null  float64
8   V8          284807 non-null  float64
9   V9          284807 non-null  float64
10  V10         284807 non-null  float64
11  V11         284807 non-null  float64
12  V12         284807 non-null  float64
13  V13         284807 non-null  float64
14  V14         284807 non-null  float64
15  V15         284807 non-null  float64
16  V16         284807 non-null  float64
17  V17         284807 non-null  float64
18  V18         284807 non-null  float64
19  V19         284807 non-null  float64
20  V20         284807 non-null  float64
21  V21         284807 non-null  float64
22  V22         284807 non-null  float64
23  V23         284807 non-null  float64
24  V24         284807 non-null  float64
25  V25         284807 non-null  float64
26  V26         284807 non-null  float64
27  V27         284807 non-null  float64
28  V28         284807 non-null  float64
29  Amount      284807 non-null  float64
30  Class       284807 non-null  int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

```
In [8]: data.isnull().sum()
```


```
Out[8]: Time      0
        V1        0
        V2        0
        V3        0
        V4        0
        V5        0
        V6        0
        V7        0
        V8        0
        V9        0
        V10       0
        V11       0
        V12       0
        V13       0
        V14       0
        V15       0
        V16       0
        V17       0
        V18       0
        V19       0
        V20       0
        V21       0
        V22       0
        V23       0
        V24       0
        V25       0
        V26       0
        V27       0
        V28       0
        Amount    0
        Class     0
        dtype: int64
```

```
In [9]: data['Class'].value_counts()
```

```
Out[9]: Class
0      284315
1         492
Name: count, dtype: int64
```

```
In [10]: data['Class'].value_counts().plot(kind = 'bar')
```

```
Out[10]: <AxesSubplot: xlabel='Class'>
```



we have a huge difference btwn fraud values or non fraud values means the data is biased so we will split the target value into two categories and then take out a sample from that and concatenate that with the data having smaller values to make a new dataset so that we can reduce the biasedness in data

```
In [11]: non_fraud = data[data.Class == 0]
         fraud = data[data.Class == 1]
```

```
In [12]: non_fraud.shape
```

```
Out[12]: (284315, 31)
```

```
In [13]: fraud.shape
```

```
Out[13]: (492, 31)
```

```
In [14]: non_fraud_sample = non_fraud.sample(600)
```

```
In [15]: non_fraud_sample
```

Out[15]:

	Time	V1	V2	V3	V4	V5	V6	
128226	78641.0	1.203948	0.131052	0.598463	0.562670	-0.616091	-0.810741	-0.1099
81009	58756.0	-2.319104	-1.556175	1.757399	-1.275513	-0.685819	0.348512	-0.7544
282522	170957.0	-2.759096	3.027481	-2.699052	-0.681087	-0.957376	-0.951570	-0.9850
67993	52784.0	1.133615	-0.086691	1.261303	1.156702	-0.983421	0.003024	-0.7217
181361	124943.0	-0.585651	1.241902	-1.488338	-0.566437	1.278687	-0.879811	1.3419
...
27839	34727.0	1.005408	-1.120883	0.431941	-0.741640	-0.933930	0.380253	-0.7914
83089	59685.0	1.278922	-0.688400	0.306704	-0.877167	-0.773223	-0.112205	-0.6844
24651	33314.0	1.317910	-1.101923	-0.663724	-1.761183	0.965597	3.552032	-1.5228
198922	132714.0	-5.203447	-7.087964	-0.061621	0.667184	3.751602	-0.096922	-0.8349
70904	54070.0	1.257282	-1.121529	-1.369897	-1.543310	1.415319	3.406259	-0.8917

600 rows × 31 columns

In [16]:

new_data = pd.concat([non_fraud_sample,fraud], axis =0)

In [17]:

new_data

Out[17]:

	Time	V1	V2	V3	V4	V5	V6	
128226	78641.0	1.203948	0.131052	0.598463	0.562670	-0.616091	-0.810741	-0.1099
81009	58756.0	-2.319104	-1.556175	1.757399	-1.275513	-0.685819	0.348512	-0.7544
282522	170957.0	-2.759096	3.027481	-2.699052	-0.681087	-0.957376	-0.951570	-0.9850
67993	52784.0	1.133615	-0.086691	1.261303	1.156702	-0.983421	0.003024	-0.7217
181361	124943.0	-0.585651	1.241902	-1.488338	-0.566437	1.278687	-0.879811	1.3419
...
279863	169142.0	-1.927883	1.125653	-4.518331	1.749293	-1.566487	-2.010494	-0.8828
280143	169347.0	1.378559	1.289381	-5.004247	1.411850	0.442581	-1.326536	-1.4137
280149	169351.0	-0.676143	1.126366	-2.213700	0.468308	-1.120541	-0.003346	-2.2347
281144	169966.0	-3.113832	0.585864	-5.399730	1.817092	-0.840618	-2.943548	-2.2080
281674	170348.0	1.991976	0.158476	-2.583441	0.408670	1.151147	-0.096695	0.2230

1092 rows × 31 columns

values of x and y

```
In [18]: x = new_data.drop(['Class'], axis = 1)
```

```
In [19]: print(x)
```

	Time	V1	V2	V3	V4	V5	V6
\							
128226	78641.0	1.203948	0.131052	0.598463	0.562670	-0.616091	-0.810741
81009	58756.0	-2.319104	-1.556175	1.757399	-1.275513	-0.685819	0.348512
282522	170957.0	-2.759096	3.027481	-2.699052	-0.681087	-0.957376	-0.951570
67993	52784.0	1.133615	-0.086691	1.261303	1.156702	-0.983421	0.003024
181361	124943.0	-0.585651	1.241902	-1.488338	-0.566437	1.278687	-0.879811
...
279863	169142.0	-1.927883	1.125653	-4.518331	1.749293	-1.566487	-2.010494
280143	169347.0	1.378559	1.289381	-5.004247	1.411850	0.442581	-1.326536
280149	169351.0	-0.676143	1.126366	-2.213700	0.468308	-1.120541	-0.003346
281144	169966.0	-3.113832	0.585864	-5.399730	1.817092	-0.840618	-2.943548
281674	170348.0	1.991976	0.158476	-2.583441	0.408670	1.151147	-0.096695
	V7	V8	V9	...	V20	V21	V22 \
128226	-0.109900	-0.011047	-0.067770	...	-0.158679	-0.205060	-0.680799
81009	-0.754404	0.950477	-0.814573	...	0.834504	0.320837	0.218636
282522	-0.985091	2.241800	-0.126389	...	0.081408	-0.228730	-0.912153
67993	-0.721792	0.273450	0.654784	...	-0.195585	0.029428	0.164201
181361	1.341948	-0.158087	0.133879	...	0.014966	0.144172	0.802194
...
279863	-0.882850	0.697211	-2.064945	...	1.252967	0.778584	-0.319189
280143	-1.413170	0.248525	-1.127396	...	0.226138	0.370612	0.028234
280149	-2.234739	1.210158	-0.652250	...	0.247968	0.751826	0.834108
281144	-2.208002	1.058733	-1.632333	...	0.306271	0.583276	-0.269209
281674	0.223050	-0.068384	0.577829	...	-0.017652	-0.164350	-0.295135
	V23	V24	V25	V26	V27	V28	Amount
128226	0.162836	0.501235	0.124477	0.068120	-0.039378	0.008476	0.89
81009	0.180391	-0.298655	0.481468	-0.293975	0.154894	-0.127781	199.73
282522	0.401224	0.626090	-0.115280	0.121893	0.124390	0.049659	9.03
67993	0.001280	0.308449	0.329003	-0.413555	0.060083	0.025442	1.00
181361	-0.340436	-1.027089	-0.203774	-0.116839	0.144424	0.054446	17.73
...
279863	0.639419	-0.294885	0.537503	0.788395	0.292680	0.147968	390.00
280143	-0.145640	-0.081049	0.521875	0.739467	0.389152	0.186637	0.76
280149	0.190944	0.032070	-0.739695	0.471111	0.385107	0.194361	77.89
281144	-0.456108	-0.183659	-0.328168	0.606116	0.884876	-0.253700	245.00
281674	-0.072173	-0.450261	0.313267	-0.289617	0.002988	-0.015309	42.53

[1092 rows x 30 columns]

```
In [20]: y = new_data["Class"]
```

```
In [21]: print(y)
```

```
128226    0
81009     0
282522    0
67993     0
181361    0
..
279863    1
280143    1
280149    1
281144    1
281674    1
Name: Class, Length: 1092, dtype: int64
```

splitting of data into training and testing sets

```
In [22]: x_train, x_test, y_train, y_test = train_test_split(x,y, random_state = 2, t
```

```
In [23]: print(x.shape, x_train.shape, x_test.shape)
```

```
(1092, 30) (982, 30) (110, 30)
```

```
In [24]: model = LogisticRegression()
```

```
In [25]: model.fit(x_train, y_train)
```

```
/home/tamanna/.local/lib/python3.11/site-packages/sklearn/linear_model/_logis
tic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
Out[25]: ▾ LogisticRegression
LogisticRegression()
```

accuracy on training data

```
In [26]: x_train_prediction = model.predict(x_train)
training_data_accuracy = accuracy_score(x_train_prediction, y_train)
```

```
In [27]: training_data_accuracy
```

```
Out[27]: 0.9439918533604889
```

accuracy on testing data


```
In [28]: x_test_prediction = model.predict(x_test)
testing_data_accuracy = accuracy_score(x_test_prediction, y_test)
```

```
In [29]: testing_data_accuracy
```

```
Out[29]: 0.9454545454545454
```

svm classifier

```
In [30]: model2 = svm.SVC(kernel = "linear")
```

```
In [31]: model2.fit(x_train, y_train)
```

```
Out[31]: SVC
SVC(kernel='linear')
```

```
In [32]: x_train_prediction = model2.predict(x_train)
```

```
In [33]: training_data_accuracy = accuracy_score(x_train_prediction, y_train)
```

```
In [34]: training_data_accuracy
```

```
Out[34]: 0.9042769857433809
```

```
In [35]: x_test_prediction = model2.predict(x_test)
```

```
In [36]: testing_data_accuracy = accuracy_score(x_test_prediction, y_test)
testing_data_accuracy
```

```
Out[36]: 0.9272727272727272
```